Arjun Subramonian

Email: arjun.subramonian@gmail.com
Website: arjunsubramonian.github.io
Publications: Semantic Scholar, Google Scholar
GitHub: github.com/ArjunSubramonian

Research Interests

AI, safety, evaluation science, governance, alignment, social impacts

Education

PhD in Computer Science; University of California, Los Angeles (2025), GPA: 3.940
Amazon Science Hub Fellow, NSF MENTOR Fellow, Eugene V. Cota-Robles Fellow
Advisors: Yizhou Sun, Kai-Wei Chang

BS in Computer Science; University of California, Los Angeles (2021), GPA: 3.927
Summa Cum Laude

Work Experience

Research Scientist; Meta FAIR AI & Society (2025)
Research Intern; Meta FAIR Cost and Alignment (2024)
Research Intern; Meta FAIR Society and Responsible AI (2022)
Research Intern; Microsoft Research FATE (2022)
Privacy Research Intern; Snap, Inc. (2021)
AllenNLP Research Engineering Intern; Allen Institute for Artificial Intelligence (2021)
Software Engineering Intern; Microsoft Corporation (2020)
Software Engineering Intern; Get Heal, Inc. (2019)

Conference Papers, Journal Papers, and Book Chapters

OpenApps: Simulating Environment Variations to Measure UI-Agent Reliability, *URL*
Karen Ullrich, …, **Arjun Subramonian**, …, Mark Ibrahim
ICLR 2026 Oral (1.13% acceptance rate), Agentic AI in the Wild @ ICLR 2026, Agents in the Wild @ ICLR 2026

Agree to Disagree? A Meta-Evaluation of LLM Misgendering, *URL*
Arjun Subramonian, Vagrant Gautam, Preethi Seshadri, Dietrich Klakow, Kai-Wei Chang, Yizhou Sun
COLM 2025 (32% acceptance rate)

An Effective Theory of Bias Amplification, *URL*
Arjun Subramonian, Samuel Bell, Levent Sagun, Elvis Dohmatob
ICLR 2025 (32.08% acceptance rate)

Strong Model Collapse, *URL*
Elvis Dohmatob, Yunzhen Feng, **Arjun Subramonian**, Julia Kempe
ICLR 2025 Spotlight (5.1% acceptance rate)

SHADES: Towards a Multilingual Assessment of Stereotypes in Large Language Models, *URL*
Margaret Mitchell, …, **Arjun Subramonian**, …, Zeerak Talat
NAACL 2025 (22% acceptance rate)

Theoretical and Empirical Insights into the Origins of Degree Bias in Graph Neural Networks, *URL*
Arjun Subramonian, Jian Kang, Yizhou Sun
NeurIPS 2024 (25.8% acceptance rate)

Understanding "Democratization" in NLP Research, *URL*
Arjun Subramonian*, Vagrant Gautam*, Dietrich Klakow, Zeerak Talat
EMNLP 2024 (20.8% acceptance rate)

Networked Inequality: Preferential Attachment Bias in Graph Neural Network Link Prediction, *URL*
Arjun Subramonian, Levent Sagun, Yizhou Sun
ICML 2024 (27.5% acceptance rate)
Graph Learning Frontiers @ NeurIPS 2023 Oral (15.4% acceptance rate)

Evaluating the Social Impact of Generative AI Systems in Systems and Society, *URL*
Irene Solaiman*, Zeerak Talat*, …, **Arjun Subramonian**
Oxford University Press Handbook on Generative AI

Motif-Driven Contrastive Learning of Graph Representations, *URL*
Shichang Zhang, Ziniu Hu, **Arjun Subramonian**, Yizhou Sun
Transactions on Knowledge and Data Engineering, Self-Supervised Learning @ WWW 2021

Factoring the Matrix of Domination: A Critical Review and Reimagination of Intersectionality in AI Fairness, *URL*
Anaelia Ovalle, **Arjun Subramonian**, Vagrant Gautam, Gilbert Gee, Kai-Wei Chang
AIES 2023 Oral (7.7% acceptance rate)

Bound by the Bounty: Collaboratively Shaping Evaluation Frameworks for Queer AI Harms, *URL*
Organizers Of QueerInAI, …, **Arjun Subramonian**, …, Jess de Jesus de Pinho Pinhal
AIES 2023 (28.9% acceptance rate)

It Takes Two to Tango: Navigating Conceptualizations of NLP Tasks and Measurements of Performance, *URL*
Arjun Subramonian, Xingdi Yuan, Hal Daumé III, Su Lin Blodgett
Findings of ACL 2023 (41.89% acceptance rate)

[Best Paper] Queer In AI: A Case Study in Community-Led Participatory AI, *URL*
Organizers of QueerInAI, …, **Arjun Subramonian**, …, Luke Stark
FAccT 2023 (24.5% acceptance rate)

On the Discrimination Risk of Mean Aggregation Feature Imputation in Graphs, *URL*
Arjun Subramonian, Kai-Wei Chang, Yizhou Sun
NeurIPS 2022 (25.6% acceptance rate)

Queer in AI, *URL*
Organizers of QueerInAI, …, **Arjun Subramonian***, …, Jeffrey Xiong*
XRDS: Crossroads, The ACM Magazine for Students, Volume 28, Issue 4

On Dyadic Fairness: Exploring and Mitigating Bias in Graph Connections, *URL*
Arjun Subramonian
ICLR 2022 Blogpost Track (32% acceptance rate)

Harms of Gender Exclusivity and Challenges in Non-Binary Representation in Language Technologies, *URL*
Sunipa Dev, Masoud Monajatipoor*, Anaelia Ovalle*, **Arjun Subramonian***, Jeff M Phillips, Kai-Wei Chang

EMNLP 2021 Oral (9.8% acceptance rate), Women in Machine Learning @ NeurIPS 2021

MOTIF-Driven Contrastive Learning of Graph Representations, *URL*
Arjun Subramonian
AAAI 2021 Undergraduate Consortium (17% acceptance rate)

Automated, Cost-Effective Optical System for Accelerated Antimicrobial Susceptibility Testing (AST) Using Deep Learning, *URL*
Calvin Brown, …, **Arjun Subramonian**, …, Aydogan Ozcan
ACS Photonics

Estimating the Ages of FGK Dwarf Stars Through the Use of GALEX FUV Magnitudes, *URL*
Sara Crandall, Graeme H. Smith, **Arjun Subramonian**, Kelly Ho, Evelyn M. Cochrane
The Astronomical Journal

Workshop Papers

Fairness Implications of GNN-to-MLP Knowledge Distillation, *URL*
Margaret Capetz, Yizhou Sun, **Arjun Subramonian**
Reliable ML @ NeurIPS 2025

Challenges to Grassroots Organization Engagement with AI Policy, *URL*
Jennifer Mickel, …, **Arjun Subramonian**
Algorithmic Collective Action @ NeurIPS 2025 Oral

Issues in Measuring the Fairness of Social Representation in Synthetic (Speech) Data, *URL*
Arjun Subramonian, Brooklyn Sheppard, Levent Sagun
Synthetic Data @ Aarhus Conference 2025

Pairwise Matching of Intermediate Representations for Fine-grained Explainability, *arXiv*
Lauren Shrack, Timm Haucke, Antoine Salaün, **Arjun Subramonian**, Sara Beery
CV4Animals @ CVPR 2025 Oral

Stop! In the Name of Flaws: Disentangling Personal Names and Sociodemographic Attributes in NLP, *URL*
Vagrant Gautam, **Arjun Subramonian**, Anne Lauscher, Os Keyes
Gender Bias in NLP @ ACL 2024

Prompting Multilingual Large Language Models to Generate Code-Mixed Texts: The Case of South East Asian Languages, *URL*
Zheng-Xin Yong, …, **Arjun Subramonian**, …, Alham Fikri Aji
Computational Approaches to Linguistic Code-Switching @ EMNLP 2023

Critical Technopolicy and Reflexivity For Proactive Gender-Inclusive NLP
Davi Liang, Anaelia Ovalle, **Arjun Subramonian**, Alicia Boyd
Queer in AI @ ACL 2023

Group Excess Risk Bound of Overparameterized Linear Regression with Constant-Stepsize SGD, *URL*
Arjun Subramonian, Levent Sagun, Kai-Wei Chang, Yizhou Sun
Trustworthy and Socially Responsible Machine Learning @ NeurIPS 2022

You Reap What You Sow: On the Challenges of Bias Evaluation Under Multilingual Settings, *URL*
Zeerak Talat, …, **Arjun Subramonian***, …, Oskar van der Wal*
Challenges & Perspectives in Creating Large Language Models @ ACL 2022

How to Make Virtual Conferences Queer-Friendly: A Guide, *URL*
Organizers of QueerInAI, …, **Arjun Subramonian**, …, Nyx McLean
Widening NLP @ EMNLP 2021

Preprints

Disentangling Geometry, Performance, and Training in Language Models, *arXiv*
Atharva Kulkarni, Jacob Mitchell Springer, **Arjun Subramonian**, Swabha Swayamdipta

When AI Benchmarks Plateau: A Systematic Study of Benchmark Saturation, *arXiv*
Mubashara Akhtar, ..., **Arjun Subramonian**, ..., Irene Solaiman

Queer NLP: A Critical Survey on Literature Gaps, Biases and Trends, *arXiv*
Sabine Weber, ..., **Arjun Subramonian**, ..., Yanan Long

auto-fpt: Automating Free Probability Theory Calculations for Machine Learning Theory, *arXiv*
Arjun Subramonian, Elvis Dohmatob

Survey of Bias In Text-to-Image Generation: Definition, Evaluation, and Mitigation, *arXiv*
Yixin Wan, **Arjun Subramonian**, …, Kai-Wei Chang

Weisfeiler and Leman Go Measurement Modeling: Probing the Validity of the WL Test, *arXiv*
Arjun Subramonian, Adina Williams, Maximilian Nickel, Yizhou Sun, Levent Sagun

BLOOM: A 176B-Parameter Open-Access Multilingual Language Model, *in submission*, *arXiv*
Teven Le Scao, ..., **Arjun Subramonian**, ..., Thomas Wolf

Fairness and Bias Mitigation: A practical guide into the AllenNLP Fairness module, *URL*
Arjun Subramonian

Public Scholarship

Queer AI Harms and Proposed Policy Interventions, *URL*
Organizers of QueerInAI, …, **Arjun Subramonian**, …

AI and Queer Communities, *URL*
Arjun Subramonian

FINDINGS: Exploring the Impact of AI, *URL*
National AI Advisory Committee, …, **Arjun Subramonian**, …

Queer in AI National AI Advisory Committee 2023 Briefing, *URL*
Organizers of QueerInAI, **Arjun Subramonian**

Rebuilding Trust: Queer in AI Approach to Artificial Intelligence Risk Management, *URL*
Organizers of QueerInAI, Ashwin S*, William Agnew*, Hetvi Jethwani*, **Arjun Subramonian***

Invited Talks

2026 - Building Queer Futures, QWER Hacks
2025 - Guest Lecture: A Critical Review and Reimagination of Intersectionality and Democratization in AI, Non-Ideal Algorithmic Fairness, Cornell Tech
2025 - An Effective Theory of Bias Amplification, ACM AI at UCLA Reading Group
2025 - Agree to Disagree? A Meta-Evaluation of LLM Misgendering, Meta FAIR; *Slides*
2025 - An Effective Theory of Bias Amplification, USC ISI Natural Language Seminar; *Event*, *Slides*, *Video*
2025 - Theoretical and Empirical Insights into Degree Bias in GNN, DEFirst Reading Group - MILA x Vector; *Video*
2025 - From "Democratization" to Personal Names: Reimagining NLP Practices Towards Justice, USC ISI Natural Language Seminar; *Event*, *Slides*, *Video*
2024 - From "Democratization" to Personal Names: Reimagining NLP Practices Towards Justice, Algorithmic Decision Systems (ADS) Reading Group, UC Berkeley; *Slides*
2024 - Queer In AI: A Case Study in Community-Led Participatory AI, Microsoft GLEAM Security; *Slides*
2024 - From "Democratization" to Personal Names: Reimagining NLP Practices Towards Justice, UC Berkeley Data Science for Social Justice Workshop; *Slides*
2024 - 'Queer-Inclusive' AI: Contesting the Power Structures and Logics of Machine Learning, IBM Research Tech for Justice Seminar; *Slides*
2024 - Queer In AI: A Case Study in Community-Led Participatory AI, Data & Society; *Slides*
2023 - Bias, Discrimination, and Power: Graphs and Natural Language, KUNGFU.AI; *Slides*
2023 - Bias and Power in NLP, USC ISI Natural Language Seminar; *Event*, *Slides*, *Video*
2022 - Bias and Power in NLP (for future consultants), Paris; *Slides*
2022 - Bias and Power in NLP, NLP Seminars at Dublin College University; *Slides*
2022 - Guest Lecture: Bias in Natural Language Processing, COM SCI 263: NLP, UCLA; *Slides*, *Video*
2022 - Queer in AI: Making AI Queer-Inclusive and Prioritizing Grassroots D&I Activism, Humlab, Umeå University; *Slides*, *Video*
2022 - Prioritizing Grassroots D&I Activism: Queer in AI, AAAI Workshop on Diversity in Artificial Intelligence; *Event*, *Slides*
2022 - Prioritizing Grassroots D&I Activism: Queer in AI, Nike Sport+AI Conference; *Event*, *Slides*
2021 - How Can I Make My Hackathon Queer-Inclusive?, Hackcon IX; *Slides*, *Video*
2019 - An Automated and Cost-Effective System for Early Antimicrobial Susceptibility Testing Using Optical Fibers and Deep Learning, UCLA HHMI Day; *Slides*

Invited Panels

2025 - Sociotechnical Perspectives & Implications Panel, Workshop on Evaluating AI in Practice, UC San Diego; *Event*
2025 - Asian Queer Experiences in Engineering: Intersectionality, Inclusion, and Institutional Change, American Society for Engineering Education
2024 - QICSE Panel, UC Berkeley
2024 - Challenges and Perspectives for queer and/or disabled communites presented by AI in HCI, ICML Queer in AI Workshop; *Event*
2024 - Artificial Intelligence and Democratic Values Roundtable, Center for AI & Digital Policy; *Event*
2024 - UCLA Computer Science PhD Visit Day Panel
2024 - UCLA ACM AI Outreach GALA Panel
2024 - Queer Computation: A Critical Unmaking, Coding Inquiry, McMaster
2024 - AI in the Time of Generative Models: Ensuring that Diverse Voices Are Heard, AAAI Workshop on Diversity in AI; *Event*
2023 - Students with Disabilities Panel, UCLA
2023 - US AI Policy: Next Steps
2023 - Queer in AI National AI Advisory Committee Briefing; *Event*, *Briefing*, *Video*, *Press*, *Findings*
2023 - Partnership on AI Inclusive Data Collection Workshop

2023 - Data & Society Salon Series on Responsible AI in Government
2022 - Students with Disabilities Panel, UCLA
2022 - Accessibility and Inclusion Panel, Neuromatch Academy; *Video*
2022 - Gender as a Variable in NLP, NAACL Queer in AI Workshop; *Event*, *Video*
2022 - NYU Center for Responsible AI Roundtable on Accessible AI Education
2022 - Co-Opting AI: Queer, NYU's Institute for Public Knowledge; *Event*, *Video*
2022 - How Do We Improve DEI in AI?, Nike Sport+AI Conference
2021 - Eye on A.I.: Equity & Inclusion in A.I. Technology, Toronto Public Library; *Video*
2021 - Intersectionality Panel, NAACL; *Video*

Honors

2024 - Two Sigma PhD Fellowship Nominee, 1 of 6 students selected by UCLA
2024 - Amazon Science Hub Fellowship; *Site*
2024 - Google PhD Fellowship Nominee, 1 of 4 students selected by UCLA
2024 - JP Morgan PhD Fellowship Nominee, 1 of 2 students selected by UCLA
2023 - LoG Conference Top-10 Reviewer
2023 - Queer in AI Ford Foundation Research Grant; *Site*
2023 - FAccT Best Paper Award; *Site*
2022 - NSF MENTOR Fellowship
2021 - Allen AI Outstanding Intern of the Year; *Site*
2021 - Major League Hacking Top 50; *Site*
2021 - UCLA Samueli School-Wide Outstanding Bachelor of Science; *Site*
2021 - UCLA Chancellor's Service Award; *Site*
2021 - UCLA Samueli Engineering Achievement Award in Student Welfare; *Site*
2021 - UCLA Eugene V. Cota-Robles Fellowship
2021 - Boeing Company Scholarship
2021 - Brian J. Lewis Endowment
2018-2021 - UCLA Dean's Honors List
2020 - Computing Research Association Outstanding Undergraduate Researcher Honorable Mention; *Site*
2020 - IBM Quantum Challenge
2019 - 3rd Place Award for Best Hack @ Rose Hack, Major League Hacking
2017 - Siemens Competition Regional Finalist; *Site*
2016 - Award of Achievement, Association for Computing Machinery, San Francisco Bay Area Professional Chapter

Service

- **Reviewing:**
  - **Conferences:** NeurIPS (2024, D&B 2024, Ethics 2024, 2023, D&B 2023, Ethics 2023), ICML (2026, 2025), ICLR (2025, Workshops 2025), FAccT (2026, 2025, 2024, 2023, 2022), LoG (2024, 2023, 2022), KDD (2024, 2023), ACL (Ethics 2024), IASEAI (2026), WACV (2026), ICCV (2025), WebConf (2024)
  - **Workshops:** NeurIPS (GLFrontiers 2023, GLFrontiers 2022, TSRML 2022), NAACL (TrustNLP 2022, WOAH 2022, SRW 2022), ACL (TrustNLP 2023, WOAH 2023, Queer in AI 2023, BigScience 2022)
- **Conference organization:** LoG Outreach Chair (2024), FAccT Socials Chair (2024), NeurIPS Affinity Workshops Chair (2022, 2023), KDD EDI Day Chair (2023), NAACL Accessibility Chair (2022)
- **Workshop and social organization:** Evaluating Evaluations: Examining Best Practices for Measuring Broader Impacts of Generative AI @ NeurIPS 2024; In Código of Hope: Possibilities of Global Solidarity and Activism in Academic Environments @ FAccT 2024; Global AI Cultures @ ICLR 2024; Collaboratively Developing

Evaluation Frameworks for Queer AI Harms @ FAccT 2022; Queer in AI: ICML (2024, 2022, 2021); NeurIPS (2021); ICLR (2024); AAAI (2021); NAACL (2025, 2022); ACL (2023)

- **Policy:** Queer in AI Policy Working Group, NIST AI Safety Institute Consortium
- **Teaching:** Computer Science 32 (2023), Computer Science Summer Institute (2022, 2023)